

Sequor — The Agentic Work Platform

Vision Brief & Analysis Walkthrough

June 2026

What this document is. A plain-language walkthrough of the Sequor platform vision and the strategic analysis behind it. It is written for a decision: *is this the right direction to build a plan on?* It is **not** a build spec, a budget, or a go-to-market plan — those come later and stay your call.

1. The idea, in one sentence

Turn the AI agent — the kind that today lives in a coder's terminal — into a single work surface where any employee, not just coders, gets their whole job done by stating what they want, while every step stays visible, controllable, and undoable.

2. The problem we're attacking

Today a staff member is the **human glue** between disconnected tools. To produce one quarterly report they hop ERP → CRM → spreadsheet → Word → an internal portal, copying data between each, holding all the context in their head.

Every task is really three things:

- an **objective** — what they want done,
- a **process** — how *this* company does it (varies company to company), and
- **data** — spread across all those tools.

The real cost isn't one slow task. It's that **the human is the integration layer** — and that doesn't scale, makes mistakes, and can't be audited.

3. The bet

Flip it: make the agent the integration layer, and the human the director.

The person says *"I want the Q3 financial report."* The agent reaches into the ERP, the CRM, and the spreadsheet, does the work, and hands back the result — without the person ever

leaving one screen. The company’s **process** becomes reusable building blocks; the **data** comes through connectors; the **objective** is just what you type.

4. Why it can win (and won’t get crushed)

This is the most important section, because there is a serious competitor.

The threat. Anthropic’s “Claude Cowork” (launched around April 2026) already does the *surface* version of this — “*an agent does your office work in one place*” — and it ships fast. If our entire pitch were “an agent does your work,” **we would already be late.**

So we don’t compete there. We win on four things underneath the surface that nobody — including Cowork — has actually delivered for non-coders:

Pillar	What it means for the user
Undo anything, see everything	Every step the agent took is recorded. You can rewind to <i>any</i> earlier step, change it, and only the affected downstream work re-runs — old versions are kept. A glass box with an undo button.
You choose the leash, up front	Before the agent runs, you pick: <i>Go ahead</i> (full autonomy), <i>Ask me once</i> , or <i>Step through with me</i> . Spending limits and approvals are built in, not bolted on.
Teams of humans <i>and</i> agents, one shared space	Several people and several agents work the same job, with everyone’s steps attributed and auditable. (Agent-to-agent teamwork is common; multiple <i>humans</i> on one shared, governed surface is the rare part.)
Share know-how safely across teams and companies	A company’s processes become reusable “recipes” you can publish, govern, and recall. This is the engine that makes the platform more valuable as more people use it.

Why competitors can’t easily copy this. The big suite vendors (Salesforce, ServiceNow, SAP, Microsoft) are locked to their own systems by business model — they can’t credibly be “works with everything.” That neutrality is open ground.

5. The honest risks (not sugar-coated)

Three things this rests on are **not proven yet**. They are flagged everywhere in the analysis so nobody sells them as fact:

1. **The “agents communicate better than humans” idea is a genuine bet.** It’s likely true for clean *handoffs*, but risky for judgment, nuance, and accountability. We treat it as something to **test cheaply**, not assume.
 2. **“Undo anything” is the strongest advantage *and* the hardest thing to build** — AI isn’t perfectly repeatable, so “re-run from step 4” is genuinely tricky, and making it legible to a non-coder is unsolved design work. Version 1 is deliberately scaled back (simple rewind, no branching).
 3. **Non-coder depth is where simple “no-code” tools historically die** at the last 20%. Our bet is that the “see everything” transparency is what makes the depth usable — but it must be *proven*, not asserted.
-

6. The good news — most of it already exists

About **80%** of the hard machinery — the governance, the trust/permission levels, the team-coordination, and the artifact-sharing engine — is **already built and shipping** across the Kailash / Terrene ecosystem. The genuinely *new* work concentrates in just three places:

- the **non-coder interface**,
- the **“undo” engine**, and
- the **safe cross-company sharing** model.

Caveat that stays attached to that 80%: existing pieces are not a finished product — it’s a head start, not a shortcut.

7. Where the current product fits

The existing **email / WhatsApp communication product becomes the first proving ground**, not a throwaway. It already quietly demonstrates several of the pillars:

- confidence-based human escalation = the “*choose the leash*” idea in miniature,
- learning from human answers = the *sharing / know-how* engine, and
- per-customer data isolation = the multi-company boundary.

It is the **lighthouse** that lets us test the risky bets on real users before betting big.

8. What we deliberately did *not* decide

Per your steer, we **did not pick a target market or first customer**. The analysis stays horizontal and proves the **capability** first; the “who do we sell to first” decision is left open, on purpose, for later.

9. Two things that need *your* decision

1. **A small correction to confirm.** Your brief defined “D/T/R” as *Decision / Task / Review*. In the governance system it actually means **Department / Team / Role** (who is accountable). The analysis uses the correct one — please confirm that matches your intent.
 2. **A premise we down-weighted.** One of your premises — *agents communicate better than humans* — is treated as a **bet to validate**, not a settled selling point. If you firmly believe it is already true, say so; it changes how hard the plans lean on it.
-

10. So what does “approve” actually mean?

Approving says: the problem, the strategy, the four pillars, the honest risks, and the “comms-as-first-proving-ground” framing are right — turn this into a build plan.

It does *not* commit you to: a budget, a timeline, building everything, or a market. Those come at the next step and remain entirely your call.

11. What’s next

The next step is **planning** (/todos): turn this analysis into a sharded, sequenced build plan. The first inputs are already identified — the cheapest, most decisive experiments first (prove the “undo” engine and the cross-system reach on the comms lighthouse) before any heavy build.

Appendix — What backs this analysis

This walkthrough is a summary. The full analysis package (~19,000 lines, red-teamed) lives in the repository under `workspaces/future-of-work/` and `specs/`:

- **Research (9 docs)** — the COC artifact system, multi-operator coordination, PACT governance, the EATP trust/posture model, the CLI-harness-as-universal-interface thesis, the transparency/versioning architecture, the competitive landscape, the work-disruption thesis, and the comms-wedge mapping.
- **Analysis (8 docs + executive summary)** — value propositions, unique selling points, the platform model, the AAA framework (Automate / Augment / Amplify), network effects,

the transparency-and-undo architecture, the 80/15/5 reuse split, and a risks/failure-point audit.

- **Plans (5)** — architecture, capability roadmap, the “undo” engine design, the trust/permission model, and the comms-wedge integration.
- **User flows (5)** — objective-to-output, rewind-and-intervene, team collaboration, recipe authoring/sharing, and the comms wedge end-to-end.
- **Specs (7 target-state platform specs)** — the authoritative description of what the platform is and does, alongside the preserved specs for the shipped comms product.
- **Red team** — independent adversarial review across four dimensions; all blocking and high-priority findings resolved. Brief coverage: 22 of 24 requirements fully addressed, 1 partially (deferred by design), 0 missing, no scope creep.

Start with `workspaces/future-of-work/01-analysis/00-EXECUTIVE-SUMMARY.md` for the deeper read.